

ASTRA

Your AI Founding Team

"You bring the idea. Astra handles everything else."

Gemini XPRIZE Hackathon 2026 · 90-Day Execution Plan · v2.0 · Confidential

6

Specialized AI agents

72hrs

Idea → live company

\$0

Technical knowledge
required

24/7

Autonomous always-on
agents

THE CONCEPT

Astra is an AI founding team for first-time startup founders.

You bring the idea. Astra handles everything else — entity formation, product development, website, market research, legal agreements, and marketing — through a stack of six autonomous specialized agents.

Most first-time founders never ship. Not because they lack a good idea — but because building a company requires six completely different skillsets simultaneously: legal, technical, design, research, marketing, and operations. Hiring for all six costs \$300,000+ per year. Astra delivers all six for less than the cost of a gym membership.

The core insight

The gap between 'I have an idea' and 'I have a live company' is not intelligence or ambition — it is execution infrastructure. First-time founders get paralyzed by the setup: forming the LLC, building the landing page, validating the market, writing the cold emails. Astra eliminates every one of those blockers autonomously, before the founder even knows they exist.

The summer timing is not accidental. Every June, millions of students finish their semester with an idea they want to build. By September, 90% of them have done nothing — not because the idea was bad, but because the first steps felt impossibly complex. Astra is built for that exact window.

THE PROBLEM

Starting a company is a full-time job before the real work even begins.

Here is what a first-time founder actually has to figure out before they can build anything:

- Which entity to form — LLC, C-Corp, S-Corp — and why it matters for future fundraising
- How to file in Delaware, what a registered agent is, and how to get an EIN from the IRS
- How to write a founder agreement so co-founders do not blow up the company later
- Whether the market is real — TAM, competitors, customer segments, pricing benchmarks
- How to build a landing page that converts without knowing how to code
- How to find first customers — cold email, ads, content, partnerships
- How to build the actual product when you are not an engineer

A wealthy founder with the right network gets help with all of this for free — advisors, connected parents, expensive accelerators. A first-generation founder figures it out alone, usually incorrectly, usually too slowly. Astra closes that gap entirely.

One text description. Six agents. A live company in 72 hours.

The founder's onboarding is a single conversation — in text. They describe their idea and Astra's agents spin up automatically, each taking ownership of their domain, each building on what the others produce. The founder never leaves the dashboard.

The Onboarding Flow

Step 1 — Founder describes the idea in text (5 minutes)

The founder types a plain-English description of their idea directly in the Astra dashboard: what it is, who it is for, and what problem it solves. No form, no template, no structured questionnaire. Just a text box. Astra extracts everything it needs from that description and distributes context to all six agents simultaneously.

○ **Hosted Llama 3.1 70B:** *Parses the founder's text description and extracts structured business context — idea summary, target customer, problem statement, and initial market hypotheses — distributed to all agents.*

○ **Gemini Flash (GCP):** *Routes the extracted context to the appropriate agents and manages orchestration across the full agent stack.*

Step 2 — Research Agent validates the market (2 hours)

While the founder does anything else, the Research Agent runs a full market analysis: total addressable market, competitor landscape, customer persona development, pricing benchmarks, and a go/no-go validation summary with honest recommendations.

○ **Hosted Llama 3.1 70B:** *Deep synthesis of research findings into a structured, readable report.*

○ **Gemini + Google Search (GCP):** *Real-time web access for current market data, competitor pricing, and recent industry news — not training cutoff data.*

Step 3 — Legal Agent forms the company (24 hours)

The Legal Agent files the Delaware LLC or C-Corp automatically via Stripe Atlas using Computer Use — navigating the actual UI, filling out the forms, completing the payment. Then retrieves the EIN from IRS.gov. Drafts the founder agreement, NDA, and IP assignment docs. Business bank account setup via Mercury is explored as a stretch goal depending on technical feasibility.

○ **Hosted Llama 3.1 70B:** *Drafts all legal documents from templates. Fills form fields. Generates the founder agreement and NDA in the company's specific context.*

○ **Antigravity Agent (GCP-hosted):** *Executes Computer Use on Stripe Atlas and IRS.gov — navigating the actual UI autonomously like a human would.*

Step 4 — Web Agent launches the landing page (48 hours)

Based on the research output and the founder's description, the Web Agent generates a complete landing page — headline, value proposition, feature highlights, waitlist form, and pricing preview. Deploys live via Framer or Vercel. The company exists on the internet before the end of Day 2.

○ **Hosted Llama 3.1 70B:** *Generates all landing page copy — headline, subtitle, feature descriptions, CTA text — in the founder's brand voice.*

○ **Google Stitch / UI Generation (GCP):** *Handles visual layout generation for the landing page structure.*

Step 5 — Marketing Agent finds the first customers (ongoing)

Builds the go-to-market plan, writes a cold email sequence targeting the ideal customer profile identified by Research, sets up Google and Meta ad campaigns within the founder's budget, and generates social content across platforms. Adjusts autonomously based on what converts.

◻ **Hosted Llama 3.1 70B:** *All copy generation — cold emails, ad copy, social posts, follow-up sequences.*

◻ **Antigravity Agent (GCP-hosted):** *Executes Computer Use on ad platforms and social accounts. Posts content. Monitors performance.*

Step 6 — Technical Agent builds features on demand

When the founder needs something built — a sign-up flow, a payment integration, a dashboard — they describe it in plain English via the agent chat interface. The Technical Agent breaks it into a spec, passes it to the hosted coding model, reviews the output, and deploys it. The founder sees the result, never the code.

◻ **Hosted DeepSeek Coder / CodeLlama:** *Generates, tests, and iterates on actual code. Runs entirely on Astra's GCP infrastructure.*

◻ **Hosted Llama 3.1 70B:** *Generates the technical spec from plain English and explains all outputs back to the founder in non-technical language.*

Step 7 — Ops Agent runs everything in the background (always-on)

Manages the founder's task list, surfaces what needs a human decision, sends the weekly progress digest, drafts routine email responses, and maintains a full persistent memory of every decision, document, and agent output across the entire company history.

◻ **Hosted Llama 3.1 70B + GCP Vector DB:** *Maintains full company memory across all sessions. Every agent output is stored and retrievable. The Ops Agent can answer 'what did we decide about pricing in Week 2?' accurately.*

One screen. Full visibility. Total control.

The dashboard is where the founder lives. They never need to open another tool — legal status, product progress, marketing performance, agent activity, and their approval queue are all in one place. Every agent reports to the dashboard. Every output flows through it.

Dashboard Layout

Dashboard Panel	What the Founder Sees	What They Can Do
Agent Status Ring	A visual ring for each of the 6 agents showing: Running / Waiting for Approval / Completed / Idle. Updated in real time as agents work.	Click any agent ring to open its full activity log and see exactly what it is working on right now.
Approval Queue	A prioritized inbox of everything that needs a human decision before the agent proceeds. Displayed as cards with full context — what the agent did, what it recommends, why it is asking.	Approve with one click, reject with a reason, or edit the output directly before approving. Agent resumes immediately.
Agent Chat Interface	A dedicated chat window per agent. The founder can open any agent and talk to it directly — ask questions, give new instructions, request a specific output, or change direction entirely.	Type naturally: 'Legal Agent — can you draft an NDA for a contractor I'm hiring?' or 'Research Agent — look into the competitor that just launched last week.'
Company Timeline	A chronological feed of every action every agent has taken — filed the LLC, launched the landing page, sent 50 cold emails, built the signup flow. The full company story in one scroll.	Filter by agent, by date, or by task type. Download any output (document, report, code) directly from the timeline.
Credit Balance & Usage	Current credit balance, monthly usage breakdown by agent, and a projection of how long credits will last at current burn rate.	Purchase additional credits inline. Upgrade plan. Set per-agent credit limits to prevent runaway usage.
Weekly Digest Preview	The Ops Agent's weekly summary — what happened, what worked, what needs attention, and the recommended priorities for the coming week.	Approve and send to co-founders, or ask the Ops Agent to revise before it goes out.

The Approval Queue — how Astra stays in control without being annoying

Agents do not ask for approval on every action — only the ones that matter. Sending a cold email to 500 people? Approval required. Updating the website headline? Approval required. Posting a social media update? Auto-approved with a notification. Filing the LLC? Approval required before payment is made. The system learns which actions the founder approves consistently and gradually reduces interruptions over time.

Agent Chat Interface

Every agent has its own chat window in the dashboard. The founder can open any agent at any time and talk to it directly — as if it were a specialist they have on retainer.

- 'Legal Agent — draft an NDA for a contractor I'm onboarding next week'
- 'Research Agent — one of our competitors just launched. Give me a full breakdown.'
- 'Marketing Agent — our open rates dropped this week. What do you think happened?'
- 'Technical Agent — I want to add a referral system to the landing page. Make it happen.'
- 'Ops Agent — what are the three most important things I should do this week?'

The agent responds with context from full company memory — it knows what was decided before, what is already in progress, and what the founder's priorities are. It is not a generic chatbot. It is a specialist that knows the company.

Six specialists. One shared brain.

Each agent owns its domain completely. Each one builds on the outputs of the others via a shared GCP vector memory store. Together they form a founding team that operates around the clock.

Agent	What It Does Autonomously	AI Model Used
 Legal Agent	Files Delaware LLC or C-Corp via Stripe Atlas (Computer Use). Obtains EIN from IRS.gov. Drafts founder agreements, NDAs, IP assignment, and vesting schedules. Monitors compliance deadlines. Mercury bank account setup is explored as a stretch goal depending on API/UI feasibility.	Hosted Llama 3.1 70B on GCP for document drafting and form field extraction. Gemini Flash (one required Google Cloud touchpoint) for orchestration routing. Antigravity agent executes Computer Use on Stripe Atlas and IRS.gov.
 Technical Agent	Takes plain-English feature descriptions. Breaks them into a technical spec. Passes to Claude Code or an open-source coding model under the hood. Opens a PR or deploys directly. Founder reviews outcomes in plain English — never touches code.	Hosted DeepSeek Coder or CodeLlama on GCP servers for code generation. Hosted Llama for spec writing and plain-English explanation. Gemini used for orchestration and routing between agents.
 Web Agent	Generates a complete landing page from brand inputs and founder's text description. Deploys live via Framer or Vercel. Builds full product website with waitlist, pricing, and auth. Iterates on design from plain-English text feedback.	Hosted Llama 3.1 for copy generation and content structure. Google's Stitch/UI generation tool for visual layout (GCP touchpoint). Framer or v0 API for deployment.
 Research Agent	Runs full market research on any idea: TAM/SAM/SOM, competitor landscape, customer personas, pricing benchmarks, and go/no-go validation summary. Updates weekly as the market shifts.	Hosted Llama 3.1 70B for deep synthesis. Google Search grounding via Gemini (GCP touchpoint) for real-time market data. Outputs structured reports and slide-ready summaries.
 Marketing Agent	Builds the go-to-market plan. Writes and sends cold email sequences. Generates social content. Sets up and manages Google Ads and Meta Ads within founder's budget. Monitors campaign performance and adjusts autonomously.	Hosted Llama for all copy generation. Antigravity agents execute Computer Use on ad platforms and social accounts. Gemini used for campaign performance analysis (GCP touchpoint).
 Ops Agent	Manages the founder's task list. Surfaces what needs a human decision. Sends weekly progress digest. Handles routine email drafts. Maintains full persistent memory of every decision, document, and agent output across the entire company history.	Hosted Llama 3.1 70B with large context for persistent company memory. All agent outputs are stored in a shared vector database on GCP. Gemini for cross-agent orchestration (GCP touchpoint).

Self-hosted models. One Google Cloud touchpoint. No API dependency.

Astra runs its own models on GCP GPU instances — not third-party API calls. This means cost control as we scale, no rate limit exposure, and no dependency on pricing decisions made by OpenAI or Anthropic. Gemini is used for the one thing it does better than anything else: Google Search grounding for real-time market data, and cross-agent orchestration within the GCP ecosystem.

Model / Service	Where It Lives	Role in Astra
Llama 3.1 70B (primary)	Self-hosted on GCP GPU instances	The workhorse. Handles document drafting, email writing, research synthesis, ops management, and all text generation tasks. Hosted on our own servers — not an API call.
DeepSeek Coder / CodeLlama	Self-hosted on GCP GPU instances	Powers the Technical Agent. Generates, reviews, and explains code in plain English. Runs on our infrastructure — no third-party coding API dependency.
Gemini Flash	Google Cloud (required GCP touchpoint)	Cross-agent orchestration, routing decisions between agents, campaign performance analysis, and tasks where Google's native integrations add clear value over hosted models.
Google Search Grounding	Google Cloud (via Gemini)	Gives the Research Agent access to real-time web data — not training cutoff data. Essential for market validation that is actually current and accurate.
Antigravity Agent Harness	Self-hosted on GCP	The persistence layer. Keeps all agents running continuously with shared state, memory, and Computer Use capabilities. One instance per active founder account.
GCP Vector Database	Google Cloud Storage + Vertex	Stores all agent outputs, company documents, and decisions in a shared embedding store. Enables the Ops Agent to maintain full company memory across sessions.
Framer / Vercel APIs	Third-party (called by Web Agent)	Web deployment targets. The Web Agent generates the site and deploys autonomously. These are execution targets, not AI models.

Why self-hosted models, not APIs?

Three reasons this is the right architectural decision:

- Cost at scale — API pricing from OpenAI or Anthropic is fine for prototypes. At 1,000 active founders running 6 agents each, the API bill becomes the business. Self-hosted inference cost drops with volume.
- No rate limits — agents run continuously, 24/7. API rate limits interrupt autonomous workflows at exactly the wrong moment. Self-hosted has no such ceiling.
- Competitive moat — a competitor can replicate a GPT-4 wrapper in a weekend. They cannot replicate a self-hosted, fine-tuned, GCP-optimized model stack with 90 days of company-specific training data without starting from scratch.

PRICING

One-time setup. Monthly agents. Credits for heavy lifting.

The Launch fee captures value from the entity formation moment. Build and Scale subscriptions capture ongoing agent usage. Credits monetize power users without pricing out early-stage founders who are bootstrapping on nothing.

LAUNCH	BUILD	SCALE
\$249 one-time setup fee	\$79/mo + credits for heavy tasks	\$149/mo more credits, priority agents
What's included: → Delaware LLC or C-Corp filing → EIN obtained automatically → Founder agreement + NDA → IP assignment docs → Landing page live on day 1 → Full market research report → Full dashboard access <i>One-time. No subscription needed to get started.</i>	Everything in Launch, plus: → All 6 agents active 24/7 → Features on demand (credits) → Marketing campaigns running → Weekly ops digest → Ongoing legal compliance → Website iterations → 20 credits/month included → Agent chat interface	Everything in Build, plus: → 50 credits/month included → Priority agent processing → Dedicated persistent memory → Investor deck generation → Pitch prep assistance → Term sheet review → Faster Computer Use queues → Early access to new agents

The Credit System — \$10 per credit

Credits power intensive tasks that require significant compute time. Build includes 20 credits per month. Scale includes 50. Additional credits purchasable at \$10 each. Unused credits roll over — no monthly expiration.

Task	Credit Cost	Real-World Equivalent
Full feature development (new screen, new endpoint)	10 credits	\$500–\$2,000 freelancer
Complete market research report (new vertical or competitor)	5 credits	\$1,500 research firm
Investor pitch deck (10 slides, designed)	8 credits	\$2,000 pitch deck designer
Full cold email campaign (500 contacts, 3-step sequence)	4 credits	\$800 marketing agency

Legal agreement (custom contract or amendment)	3 credits	\$400 per lawyer hour
Website redesign or full visual refresh	6 credits	\$3,000 web designer
Google/Meta ad campaign setup and launch	4 credits	\$1,000 marketing consultant

The value comparison

A founder using Astra's full agent stack for one month gets: entity formation (\$500+ lawyer), market research (\$1,500 firm), landing page (\$3,000 designer), cold email campaign (\$800 agency), one feature built (\$1,500 freelancer). That is \$7,300 in equivalent professional services. Launch + Build = \$328. The ROI conversation is effortless.

From nothing on Day 1 to \$1,000 MRR by August 17.

Product and revenue run in parallel from Day 1. The XPRIZE deadline is the forcing function — every week has a clear milestone for both tracks.

Week	Product	Revenue
1–2	Scaffold the app and GCP infrastructure. Self-host Llama 3.1 70B and DeepSeek Coder on GCP GPU instances. Build Legal Agent Computer Use flow (Stripe Atlas + IRS.gov). Build Research Agent with Google Search grounding via Gemini. Wire Gemini Flash as the cross-agent orchestrator. Founder text intake endpoint live.	Set up Stripe (Launch \$249, Build \$79/mo, Scale \$149/mo, credits \$10 each). Identify 30 target founders: summer hackathon groups, university Discord servers, Indie Hackers, X/Twitter 'building in public' community.
3–4	Build Web Agent (Framer deploy). Build Marketing Agent (cold email + social copy). Build Ops Agent with persistent GCP vector memory. Wire all agents into shared context. Build founder dashboard: agent status rings, approval queue, company timeline, credit balance. Agent chat interface live.	Soft launch to 10 free beta founders. Full Launch package at no cost in exchange for being a case study. Run their companies through the full stack. Record every friction point. Run Astra's own Research Agent on Astra — validate the market.
5–6	Full approval flow live — every agent output flows into the approval queue before executing. Fix top 3 beta friction points. Build onboarding flow: founder types their idea, agents spin up, LLC filing begins within 24 hours, landing page live within 48 hours.	Flip 3 beta founders to paid Launch (\$249 × 3 = \$747). Launch Build plan (\$79/mo). Cold email 50 summer founders: 'From idea to live company in 72 hours.' Post first case study. Revenue checkpoint: \$747+ one-time + first recurring.
7–8	Harden all agents. Add investor deck generation to Technical Agent. Add pitch coaching to Ops Agent. Begin exploring Mercury bank account setup feasibility — test UI/API, go/no-go decision by end of Week 8. Referral mechanic: refer 1 founder = \$25 credit.	Press outreach: Product Hunt, The Hustle, Indie Hackers newsletter. Partner with 2 university entrepreneurship centers. Target: 8 paying customers. Revenue checkpoint: \$1,500+ combined.
9–10	Launch Scale plan (\$149/mo). Build cohort access for university accelerators. Feature freeze begins on core agents — stability over new features. XPRIZE KPI documentation starts.	University accelerator deal: \$1,500–\$3,000 per cohort of 10 founders. First institutional deal target. Individual subscriptions growing. Revenue checkpoint: \$2,500+ combined.
11–13	Final hardening. XPRIZE submission prep. Astra's own company run entirely by Astra — every agent active, all outputs documented. Build the submission demo: live walkthrough of a founder going from text description to live company in real time.	Final push: 10+ paying customers across all tiers. XPRIZE submission. Revenue checkpoint: \$1,000+ MRR confirmed.

Day 90 target: \$1,000+ MRR. 10+ paying customers. All 6 agents live. XPRIZE submitted. Astra's own company run by Astra.

Five moats. Each one growing every week.

Astra is not a collection of integrations. It is an orchestration layer with self-hosted models, compound memory, and autonomous execution that gets stronger with every founder onboarded.

Full-Stack Agent Orchestration Every competitor does one thing — Stripe Atlas does legal, Framer does websites, Mailchimp does email. Astra orchestrates all of them through a single founder interface with shared context and memory. The agents build on each other's outputs.
Self-Hosted Models — No API Dependency Running Llama and DeepSeek on our own GCP servers means we are not subject to third-party API rate limits, pricing changes, or downtime. As usage scales, our cost per inference drops. Competitors relying on OpenAI or Anthropic APIs face unpredictable costs and usage caps.
Compound Company Memory The Ops Agent maintains a persistent vector memory of the entire company — every decision, document, campaign, and feature. After 90 days, Astra knows a founder's company better than any human consultant. This memory is locked to the platform — it cannot be exported to a competitor.
Computer Use for Legal Execution Filing an LLC autonomously via Computer Use on Stripe Atlas is a capability no existing tool has. It is not an API — it is an agent performing the actual human workflow in real time. Competitors cannot replicate this without persistent agent infrastructure.
The Approval Queue as Trust Layer Most AI tools either require too much human input or go fully autonomous and make mistakes. Astra's approval queue threads the needle — agents do all the work, humans make the final call on anything consequential. This builds trust without adding friction.

XPRIZE SUBMISSION STRATEGY

The story judges haven't heard before.

A self-hosted AI founding team that autonomously forms legal entities, builds products, and acquires customers — with a live portfolio of companies it helped create in 90 days.

Business Viability KPIs

Monthly Recurring Revenue (Build + Scale subscriptions)	\$1,000+
Launch one-time fees collected	\$2,490+
Total paying customers	10+
30-day customer retention	80%+

Net Promoter Score	8+ / 10
Gross margin (after GCP hosting costs)	65%+
University accelerator partnerships signed	2+

AI-Native Operations KPIs

LLCs / C-Corps filed autonomously via Computer Use	10+
Time from text description to live landing page	< 48 hours
Market research reports generated	15+
Features built via Technical Agent on demand	20+
Cold email sequences launched autonomously	10+
Agent tasks completed with zero human intervention	500+
Items processed through the approval queue	200+
Founder hours saved vs. doing everything manually	1,000+

Submission Narrative — Five Acts

Act I — The Inequality

Wealthy founders get advisors, lawyers, technical cofounders, and accelerator access. First-generation founders get a blank Google Doc. The gap is not ideas — it is infrastructure. Astra closes it.

Act II — The 72-Hour Demo

A founder with no technical background types their idea into a text box. 72 hours later: a Delaware LLC filed, a live landing page, a validated market research report, and the first cold email campaign running. Show the receipts — real filings, real URLs, real emails sent.

Act III — The Architecture

Self-hosted Llama and DeepSeek on GCP. Gemini for Google Search grounding and orchestration. Antigravity agents for autonomous Computer Use. A shared vector memory that makes six agents act like one team. This is not a wrapper — it is a ground-up agentic system.

Act IV — The Portfolio

The companies Astra helped create in 90 days. Real founders, real entities, real landing pages, real first customers. Each one is proof. Each founder's story is a testimonial.

Act V — The Scale

10 founders in 90 days is the proof of concept. 1,000 per cohort is the university accelerator model. Every new founder makes Astra's models smarter for the next one. Show the flywheel.

The meta-proof move

Astra's XPRIZE submission is prepared by Astra. The market research was run by the Research Agent. The submission deck was built by the Technical Agent. The outreach was drafted by the Marketing Agent. The legal docs are filed. The landing page is live. Astra built Astra. No competitor can replicate this.

RISK REGISTER

Honest risks. Real mitigations.

Risk	Level	Mitigation
Self-hosted model performance inconsistent	HIGH	Start with Llama 3.1 70B — well-benchmarked for instruction following and document generation. Run evals on all agent tasks in Week 1 before building on top. Fall back to larger model if quality is insufficient.
GCP GPU hosting costs spike	HIGH	Set hard per-founder compute budgets enforced at the infrastructure layer. Batch non-urgent agent tasks. Use quantized models (4-bit) where quality allows. Monitor burn rate weekly against per-customer COGS target.
Stripe Atlas Computer Use breaks	HIGH	Build direct Delaware registered agent API integration as fallback (Harvard Business Services, InCorp). Computer Use is primary — API fallback is insurance, not the plan.
Legal liability from document generation	HIGH	Explicit disclaimer on all documents: 'AI-generated document preparation — not legal advice. Review with a licensed attorney before signing.' Do not use the word 'legal advice' anywhere in the product.
Mercury bank account not feasible in time	MED	Flag this as a stretch goal from Day 1. Test feasibility in Week 7–8. If UI automation is too fragile, remove from pitch entirely. Entity formation works without it — do not let it block the core flow.
Founders churn after Launch fee	MED	Design Launch package to leave clear gaps that Build fills. LLC is done but marketing, coding, and ops agents only activate on Build. The setup creates appetite — subscription delivers.
Too broad — agents are shallow across 6 functions	HIGH	Nail Legal + Research + Web in Phase 1. Add Technical + Marketing in Phase 2. Ops Agent last. Ship 3 great agents before 6 mediocre ones.

Day 1 Checklist.

Two tracks. Both start today. Assign before anyone opens a laptop.

Engineering — Morning

- Register astra.ai, withastra.com, or tryastra.co — domain first
- Create GCP project — enable Vertex AI, Cloud Storage, GPU quota for model hosting
- Request GCP GPU quota increase immediately — this can take 24–48 hours to approve
- Scaffold Next.js app, push to GitHub, deploy to Vercel

Engineering — Afternoon

- Spin up first GCP GPU instance. Deploy Llama 3.1 70B via Ollama or vLLM. Test a document generation prompt — make sure quality is acceptable before building on top of it
- Wire first Computer Use flow: get the Antigravity agent to open Stripe Atlas and navigate to the LLC filing form. First milestone: fill one field autonomously
- Set up GCP vector store for agent memory — every agent output will write here
- Assign agent owners: one team member per agent. No one owns two

Business & Research — Morning

- Set up Stripe: Launch (\$249 one-time), Build (\$79/mo), Scale (\$149/mo), credits (\$10 each)
- Write the legal disclaimer copy for all document outputs — get one attorney to review before launch
- Find 30 target founders: search Twitter/X 'building in public summer 2026', Indie Hackers, YC Startup School Discord
- Join every relevant summer founder community — your market is already assembled online

Business & Research — Afternoon

- Run Astra's Research Agent on Astra itself — who are the real competitors (Stripe Atlas, Clerky, Doola, Notion AI), what do they charge, what do they miss
- Cold email 10 summer founders today: 'You have an idea. We built an AI founding team. 72 hours from text description to live company. Free for 5 beta users.'
- Post Day 1 build log on X, LinkedIn, and relevant Discord servers. Let Astra market Astra from Day 1
- Set up KPI tracker: MRR / customers / LLCs filed / features built / agent tasks / approval queue items. Update every Sunday